

# Implementación de CAMELON

CAMELON: A System for Crime Metadata Extraction and Spatiotemporal Visualization From Online News Articles

<https://ieeexplore.ieee.org/document/10424974>

Esta implementación solo cubre la parte de crime-classification

```
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

!unzip -q /content/drive/MyDrive/CIENCIA\ DE\
DATOS/tweets_predict10k_2020.zip

import pandas as pd
ruta = "/content/tweets_predict10k_2020.csv"
df = pd.read_csv(ruta)

<ipython-input-1-976793c0ea66>:3: DtypeWarning: Columns (43) have
mixed types. Specify dtype option on import or set low_memory=False.
    df = pd.read_csv(ruta)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101630 entries, 0 to 101629
Data columns (total 44 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0.1                          101630 non-null int64
1   Unnamed: 0                            101630 non-null int64
2   ID                                    101630 non-null int64
3   url                                  101630 non-null object
4   Date                                101630 non-null object
5   lang                                101630 non-null object
6   Description                          101630 non-null object
7   replyCount                          101630 non-null int64
8   retweetCount                         101630 non-null int64
9   likeCount                           101630 non-null int64
10  quoteCount                           101630 non-null int64
11  bookmarkedCount                      101630 non-null int64
12  conversationId                       101630 non-null int64
13  hashtags                             101630 non-null object
14  cashtags                             101630 non-null object
15  mentionedUsers                       101630 non-null object
16  links                                101630 non-null object
```

```

17 source 101630 non-null object
18 sourceUrl 101630 non-null object
19 sourceLabel 101630 non-null object
20 possibly_sensitive 46600 non-null object
21 _type 101630 non-null object
22 user_id 101630 non-null int64
23 user_username 101630 non-null object
24 user_followersCount 101630 non-null int64
25 inReplyToUser_username 30204 non-null object
26 inReplyToUser_displayname 30199 non-null object
27 inReplyToUser__type 30205 non-null object
28 Year 101630 non-null int64
29 Month 101630 non-null int64
30 dayOfWeek 101630 non-null int64
31 dayOfMonth 101630 non-null int64
32 dayOfYear 101630 non-null int64
33 weekOfMonth 101630 non-null int64
34 weekOfYear 101630 non-null int64
35 Hour 101630 non-null int64
36 Minute 101630 non-null int64
37 Hour_Zone 101630 non-null object
38 BusinessHour 101630 non-null int64
39 Weekend 101630 non-null int64
40 Season 101630 non-null object
41 Holiday 101630 non-null int64
42 CleanDescription 101578 non-null object
43 Category 15000 non-null object
dtypes: int64(23), object(21)
memory usage: 34.1+ MB

df["Category"].unique()

array(['SOCIAL_COMMENTARY', 'NON_CRIME_RELATED', 'KIDNAPPING',
      'HOMICIDE',
      'ASSAULT', 'NEWS_MEDIA_MENTION', 'THEFT', 'BURGLARY',
      'BATTERY',
      'WEAPONS VIOLATION', 'OTHER OFFENSE', 'OFFENSE INVOLVING
CHILDREN',
      'PUBLIC PEACE VIOLATION', 'INTIMIDATION', 'ARSON', 'AMBIGUOUS',
      'CRIMINAL DAMAGE', 'CRIMINAL SEXUAL ASSAULT', 'ROBBERY',
      'NARCOTICS', 'INTERFERENCE WITH PUBLIC OFFICER',
      'FICTIONAL_CONTENT', 'MOTOR VEHICLE THEFT', 'DECEPTIVE
PRACTICE',
      'CRIMINAL TRESPASS', 'SEX OFFENSE', 'LIQUOR LAW VIOLATION',
      'CONCEALED CARRY LICENSE VIOLATION', 'THREAT',
      'CRIM SEXUAL ASSAULT', 'GAMBLING', 'HUMAN TRAFFICKING',
      'PROSTITUTION', 'HUMOR_OR_SATIRE', nan], dtype=object)

# Lista de clases irrelevantes o no relacionadas a crimen
non_crime_classes = [

```

```

    "FICTIONAL_CONTENT",
    "SOCIAL_COMMENTARY",
    "NEWS_MEDIA_MENTION",
    "HUMOR_OR_SATIRE",
    "NON_CRIME_RELATED",
    "AMBIGUOUS"
]

# Reemplazar esas clases por una sola etiqueta: "NON_CRIME"
df["Category"] = df["Category"].replace(non_crime_classes,
"NON_CRIME")

# Mapeo de categorías redundantes a categorías simplificadas
category_mapping = {
    "CRIMINAL SEXUAL ASSAULT": "SEXUAL ASSAULT",
    "CRIM SEXUAL ASSAULT": "SEXUAL ASSAULT",
    "SEX OFFENSE": "SEXUAL ASSAULT",

    "NARCOTICS": "DRUG OFFENSE",
    "OTHER NARCOTIC VIOLATION": "DRUG OFFENSE",

    "BATTERY": "ASSAULT",
    "ASSAULT": "ASSAULT",

    "BURGLARY": "THEFT",
    "ROBBERY": "THEFT",
    "THEFT": "THEFT",
    "MOTOR VEHICLE THEFT": "THEFT",

    "LIQUOR LAW VIOLATION": "WEAPONS/LIQUOR VIOLATION",
    "CONCEALED CARRY LICENSE VIOLATION": "WEAPONS/LIQUOR VIOLATION",
    "WEAPONS VIOLATION": "WEAPONS/LIQUOR VIOLATION",

    "PUBLIC PEACE VIOLATION": "PUBLIC ORDER OFFENSE",
    "INTERFERENCE WITH PUBLIC OFFICER": "PUBLIC ORDER OFFENSE",
    "INTIMIDATION": "PUBLIC ORDER OFFENSE",

    "OBSCENITY": "INDECENCY",
    "PUBLIC INDECENCY": "INDECENCY"
}

# Aplicar el mapeo
df["Category"] = df["Category"].replace(category_mapping)

df["Category"].unique()

array(['NON_CRIME', 'KIDNAPPING', 'HOMICIDE', 'ASSAULT', 'THEFT',
      'WEAPONS/LIQUOR VIOLATION', 'OTHER OFFENSE',
      'OFFENSE INVOLVING CHILDREN', 'PUBLIC ORDER OFFENSE', 'ARSON',
      'CRIMINAL DAMAGE', 'SEXUAL ASSAULT', 'DRUG OFFENSE',
      'DECEPTIVE PRACTICE', 'CRIMINAL TRESPASS', 'THREAT',

```

```
'GAMBLING',  
      'HUMAN TRAFFICKING', 'PROSTITUTION', nan], dtype=object)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 101630 entries, 0 to 101629
```

```
Data columns (total 44 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0.1	101630 non-null	int64
1	Unnamed: 0	101630 non-null	int64
2	ID	101630 non-null	int64
3	url	101630 non-null	object
4	Date	101630 non-null	object
5	lang	101630 non-null	object
6	Description	101630 non-null	object
7	replyCount	101630 non-null	int64
8	retweetCount	101630 non-null	int64
9	likeCount	101630 non-null	int64
10	quoteCount	101630 non-null	int64
11	bookmarkedCount	101630 non-null	int64
12	conversationId	101630 non-null	int64
13	hashtags	101630 non-null	object
14	cashtags	101630 non-null	object
15	mentionedUsers	101630 non-null	object
16	links	101630 non-null	object
17	source	101630 non-null	object
18	sourceUrl	101630 non-null	object
19	sourceLabel	101630 non-null	object
20	possibly_sensitive	46600 non-null	object
21	_type	101630 non-null	object
22	user_id	101630 non-null	int64
23	user_username	101630 non-null	object
24	user_followersCount	101630 non-null	int64
25	inReplyToUser_username	30204 non-null	object
26	inReplyToUser_displayname	30199 non-null	object
27	inReplyToUser__type	30205 non-null	object
28	Year	101630 non-null	int64
29	Month	101630 non-null	int64
30	dayOfWeek	101630 non-null	int64
31	dayOfMonth	101630 non-null	int64
32	dayOfYear	101630 non-null	int64
33	weekOfMonth	101630 non-null	int64
34	weekOfYear	101630 non-null	int64
35	Hour	101630 non-null	int64
36	Minute	101630 non-null	int64
37	Hour_Zone	101630 non-null	object
38	BusinessHour	101630 non-null	int64
39	Weekend	101630 non-null	int64

40	Season	101630	non-null	object
41	Holiday	101630	non-null	int64
42	CleanDescription	101578	non-null	object
43	Category	15000	non-null	object

dtypes: int64(23), object(21)

memory usage: 34.1+ MB

```
df["Category"].value_counts()
```

Category	
NON_CRIME	5417
HOMICIDE	4658
ASSAULT	2291
PUBLIC ORDER OFFENSE	782
WEAPONS/LIQUOR VIOLATION	560
OTHER OFFENSE	359
THEFT	232
OFFENSE INVOLVING CHILDREN	184
DRUG OFFENSE	166
KIDNAPPING	118
ARSON	110
CRIMINAL DAMAGE	39
SEXUAL ASSAULT	39
DECEPTIVE PRACTICE	18
HUMAN TRAFFICKING	13
PROSTITUTION	6
GAMBLING	4
CRIMINAL TRESPASS	2
THREAT	2

Name: count, dtype: int64

```
import pandas as pd
import torch
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from datasets import Dataset
from transformers import (
    XLMLRobertaTokenizerFast,
    XLMLRobertaForSequenceClassification,
    Trainer,
    TrainingArguments,
)
from torch.nn import CrossEntropyLoss

# Filtrar filas con categoría válida
df_filtered = df[df['Category'].notna()].copy()
df_filtered = df_filtered.dropna(subset=['CleanDescription'])

# Codificar las categorías a etiquetas numéricas
le = LabelEncoder()
```

```

df_filtered['label'] = le.fit_transform(df_filtered['Category'])
num_labels = len(le.classes_)

class_counts = df_filtered['label'].value_counts().sort_index()
weights = 1.0 / class_counts
weights = weights / weights.sum() # normalizar pesos
weights = torch.tensor(weights.values, dtype=torch.float).to('cuda' if
torch.cuda.is_available() else 'cpu')

# --- Filtrar clases con pocas muestras ---
min_samples = 100
class_counts = df_filtered['label'].value_counts()
valid_labels = class_counts[class_counts >= min_samples].index
df_filtered = df_filtered[df_filtered['label'].isin(valid_labels)]

# --- Reindexar etiquetas ---
le = LabelEncoder()
df_filtered['label'] = le.fit_transform(df_filtered['label'])

# --- División en entrenamiento y prueba ---
train_df, test_df = train_test_split(
    df_filtered[['CleanDescription', 'label']],
    test_size=0.2,
    stratify=df_filtered['label'],
    random_state=42
)

<ipython-input-12-14f84e9daea5>:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    df_filtered['label'] = le.fit_transform(df_filtered['label'])

tokenizer = XLNetTokenizerFast.from_pretrained("xlm-roberta-
base")

# --- Tokenización segura ---
def tokenize_function(examples):
    return tokenizer(
        examples["CleanDescription"],
        truncation=True,
        padding='max_length',
        max_length=128
    )

/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/
_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.

```

To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret in your Google Colab and restart your session. You will be able to reuse this secret in all of your notebooks. Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(
```

```
# --- Eliminar índice heredado (muy importante) ---
```

```
train_df = train_df.reset_index(drop=True)
```

```
test_df = test_df.reset_index(drop=True)
```

```
# --- Conversión a HuggingFace Datasets ---
```

```
train_dataset = Dataset.from_pandas(train_df)
```

```
test_dataset = Dataset.from_pandas(test_df)
```

```
tokenized_train = train_dataset.map(tokenize_function, batched=True)
```

```
tokenized_test = test_dataset.map(tokenize_function, batched=True)
```

```
{"model_id": "651755be92dc4c7d8473144548c45151", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "da07a4031e3c4e3991567d3d5dc91cb8", "version_major": 2, "version_minor": 0}
```

```
model = XLMRobertaForSequenceClassification.from_pretrained("xlm-roberta-base", num_labels=num_labels)
```

Some weights of XLMRobertaForSequenceClassification were not initialized from the model checkpoint at xlm-roberta-base and are newly initialized: ['classifier.dense.bias', 'classifier.dense.weight', 'classifier.out\_proj.bias', 'classifier.out\_proj.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
class WeightedLossTrainer(Trainer):
```

```
    def __init__(self, class_weights=None, *args, **kwargs):
```

```
        self.class_weights = class_weights
```

```
        super().__init__(*args, **kwargs)
```

```
    def compute_loss(self, model, inputs, return_outputs=False, **kwargs): # <- ¡Aquí el cambio!
```

```
        labels = inputs.get("labels")
```

```
        outputs = model(**inputs)
```

```
        logits = outputs.logits
```

```
        loss_fct = CrossEntropyLoss(weight=self.class_weights)
```

```
        loss = loss_fct(logits.view(-1, model.config.num_labels), labels.view(-1))
```

```

        return (loss, outputs) if return_outputs else loss
from sklearn.metrics import accuracy_score, f1_score, precision_score,
recall_score

def compute_metrics(eval_pred):
    logits, labels = eval_pred
    predictions = logits.argmax(axis=-1)

    return {
        "accuracy": accuracy_score(labels, predictions),
        "f1": f1_score(labels, predictions, average="weighted"),
        "precision": precision_score(labels, predictions,
average="weighted"),
        "recall": recall_score(labels, predictions,
average="weighted"),
    }

training_args = TrainingArguments(
    output_dir="./results",
    eval_strategy="epoch",
    per_device_train_batch_size=64,
    per_device_eval_batch_size=64,
    num_train_epochs=20,
    save_strategy="epoch",
    logging_dir="./logs",
    logging_steps=10,
    report_to="none",
    load_best_model_at_end=True,
)

trainer = WeightedLossTrainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_train,
    eval_dataset=tokenized_test,
    tokenizer=tokenizer,
    class_weights=weights,
    compute_metrics=compute_metrics
)

trainer.train()

<ipython-input-16-76e60c31d594>:4: FutureWarning: `tokenizer` is
deprecated and will be removed in version 5.0.0 for
`WeightedLossTrainer.__init__`. Use `processing_class` instead.
  super().__init__(*args, **kwargs)

<IPython.core.display.HTML object>

```



```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_
_classification.py:1565: UndefinedMetricWarning: Precision is ill-
defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification
.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification
.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification
.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification
.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification
.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification
.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification
.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))

```

```
set to 0.0 in labels with no predicted samples. Use `zero_division`  
parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being  
set to 0.0 in labels with no predicted samples. Use `zero_division`  
parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being  
set to 0.0 in labels with no predicted samples. Use `zero_division`  
parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being  
set to 0.0 in labels with no predicted samples. Use `zero_division`  
parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being  
set to 0.0 in labels with no predicted samples. Use `zero_division`  
parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being  
set to 0.0 in labels with no predicted samples. Use `zero_division`  
parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being  
set to 0.0 in labels with no predicted samples. Use `zero_division`  
parameter to control this behavior.
```

```

    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
Could not locate the best model at
./results/checkpoint-930/pytorch_model.bin, if you are running a
distributed training on multiple nodes, you should activate `--
save_on_each_node`.

TrainOutput(global_step=3720, training_loss=0.10442525179636093,
metrics={'train_runtime': 6442.183, 'train_samples_per_second':
36.935, 'train_steps_per_second': 0.577, 'total_flos':
1.565355006091776e+16, 'train_loss': 0.10442525179636093, 'epoch':
20.0})

model_path = "/content/drive/MyDrive/CIENCIA DE DATOS/modelos-xmlr"

# Guardar el modelo y el tokenizer en la ruta especificada
trainer.save_model(model_path)
tokenizer.save_pretrained(model_path)

('/content/drive/MyDrive/CIENCIA DE
DATOS/modelos-xmlr/tokenizer_config.json',
'/content/drive/MyDrive/CIENCIA DE
DATOS/modelos-xmlr/special_tokens_map.json',
'/content/drive/MyDrive/CIENCIA DE
DATOS/modelos-xmlr/sentencepiece.bpe.model',
'/content/drive/MyDrive/CIENCIA DE
DATOS/modelos-xmlr/added_tokens.json',
'/content/drive/MyDrive/CIENCIA DE
DATOS/modelos-xmlr/tokenizer.json')

trainer.evaluate()

```

<IPython.core.display.HTML object>

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))

{'eval_loss': 0.6859170794487,
 'eval_accuracy': 0.6292436974789916,
 'eval_f1': 0.5548542189524683,
 'eval_precision': 0.5041606520314937,
 'eval_recall': 0.6292436974789916,
 'eval_runtime': 17.0244,
 'eval_samples_per_second': 174.749,
 'eval_steps_per_second': 2.761,
 'epoch': 20.0}
```